

Lab 1 - Protocol Layers

Step 1: Capture a Trace

Proceed as follows to capture a trace of network traffic; alternatively, you may use a supplied trace. We want this trace to look at the protocol structure of packets. A simple Web fetch of a URL from a server of your choice to your computer, which is the client, will serve as traffic.

1. Pick a URL and fetch it with wget or curl. For example, “wget http://www.google.com” or “curl http://www.google.com”. This will fetch the resource and either write it to a file (wget) or to the screen (curl). You are checking to see that the fetch works and retrieves some content. A successful example is shown below (with added highlighting) for wget. You want a single response with status code “200 OK”. If the fetch does not work then try a different URL; if no URLs seem to work then debug your use of wget/curl or your Internet connectivity.

```
wsl@DESKTOP-OLHVFBO:/mnt/c/Users/esam5/Desktop/School/Fall 2020/CSCI-415/Labs/Lab-1/Data-Communications-and-Networking$ wget "google.com"
Will not apply HSTS. The HSTS database must be a regular and non-world-writable file.
ERROR: could not open HSTS store at '/home/wsl/.wget-hsts'. HSTS will be disabled.
--2020-09-30 16:06:57-- http://google.com/
Resolving google.com (google.com)... 172.217.10.46, 2607:f8b0:4006:814::200e
Connecting to google.com (google.com)|172.217.10.46|:80... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: http://www.google.com/ [following]
--2020-09-30 16:06:57-- http://www.google.com/
Resolving www.google.com (www.google.com)... 172.217.9.228, 2607:f8b0:4006:813::2004
Connecting to www.google.com (www.google.com)|172.217.9.228|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: unspecified [text/html]
Saving to: 'index.html'

index.html           [ <=> ] 13.71
K --.-KB/s    in 0s

2020-09-30 16:06:57 (45.4 MB/s) - 'index.html' saved [14034]
```

2. *Close unnecessary browser tabs and windows.* By minimizing browser activity you will stop your computer from fetching unnecessary web content, and avoid incidental traffic in the trace.

3. *Launch Wireshark and start a capture with a filter of “tcp port 80”* and check “enable net-work name resolution”. This filter will record only standard web traffic and not other kinds of packets that your computer may send. The checking will translate the addresses of the computers sending and receiving packets into names, which should help you to recognize whether the packets are going to or from your computer.

4. *When the capture is started, repeat the web fetch using wget/curl above.* This time, the packets will be recorded by Wireshark as the content is transferred.

5. After the fetch is successful, return to Wireshark and use the menus or buttons to stop the trace. If you have succeeded, the upper Wireshark window will show multiple packets, and most likely it will be full. How many packets are captured will depend on the size of the web page, but there should be at least 8 packets in the trace, and typically 20-100, and many of these packets will be colored green. An example is shown below. Congratulations, you have captured a trace!

1	0.000000	192.168.1.151	172.217.6.206	TCP	66 6494 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 WS=256 SACK_PERM=1
2	0.005111	172.217.6.206	192.168.1.151	TCP	66 80 → 6494 [SYN, ACK] Seq=0 Ack=1 Win=65535 Len=0 MSS=1430 SACK_PERM=1 WS=256
3	0.005212	192.168.1.151	172.217.6.206	TCP	54 6494 → 80 [ACK] Seq=1 Ack=1 Win=262912 Len=0
4	0.025028	192.168.1.151	172.217.6.206	HTTP	191 GET / HTTP/1.1
5	0.030344	172.217.6.206	192.168.1.151	TCP	60 80 → 6494 [ACK] Seq=1 Ack=138 Win=66816 Len=0
6	0.039274	172.217.6.206	192.168.1.151	HTTP	582 HTTP/1.1 301 Moved Permanently (text/html)
7	0.061316	192.168.1.151	172.217.10.36	TCP	66 6495 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 WS=256 SACK_PERM=1
8	0.065974	172.217.10.36	192.168.1.151	TCP	66 80 → 6495 [SYN, ACK] Seq=0 Ack=1 Win=60720 Len=0 MSS=1380 SACK_PERM=1 WS=256
9	0.066144	192.168.1.151	172.217.10.36	TCP	54 6495 → 80 [ACK] Seq=1 Ack=1 Win=262144 Len=0
10	0.080132	192.168.1.151	172.217.6.206	TCP	54 6494 → 80 [ACK] Seq=138 Ack=529 Win=262400 Len=0
11	0.085384	192.168.1.151	172.217.10.36	HTTP	195 GET / HTTP/1.1
12	0.089871	172.217.10.36	192.168.1.151	TCP	60 80 → 6495 [ACK] Seq=1 Ack=142 Win=61952 Len=0
13	0.142500	172.217.10.36	192.168.1.151	TCP	1484 80 → 6495 [ACK] Seq=1 Ack=142 Win=61952 Len=1430 [TCP segment of a reassembled PDU]
14	0.142500	172.217.10.36	192.168.1.151	TCP	1484 80 → 6495 [ACK] Seq=1431 Ack=142 Win=61952 Len=1430 [TCP segment of a reassembled PDU]
15	0.142500	172.217.10.36	192.168.1.151	TCP	1484 80 → 6495 [ACK] Seq=2861 Ack=142 Win=61952 Len=1430 [TCP segment of a reassembled PDU]
16	0.142500	172.217.10.36	192.168.1.151	TCP	1484 80 → 6495 [ACK] Seq=4291 Ack=142 Win=61952 Len=1430 [TCP segment of a reassembled PDU]
17	0.142500	172.217.10.36	192.168.1.151	TCP	1484 80 → 6495 [ACK] Seq=5721 Ack=142 Win=61952 Len=1430 [TCP segment of a reassembled PDU]
18	0.142500	172.217.10.36	192.168.1.151	TCP	1484 80 → 6495 [ACK] Seq=7151 Ack=142 Win=61952 Len=1430 [TCP segment of a reassembled PDU]
19	0.142500	172.217.10.36	192.168.1.151	TCP	1484 80 → 6495 [ACK] Seq=8581 Ack=142 Win=61952 Len=1430 [TCP segment of a reassembled PDU]
20	0.142500	172.217.10.36	192.168.1.151	TCP	1484 80 → 6495 [ACK] Seq=10011 Ack=142 Win=61952 Len=1430 [TCP segment of a reassembled PDU]
21	0.142500	172.217.10.36	192.168.1.151	TCP	1484 80 → 6495 [ACK] Seq=11441 Ack=142 Win=61952 Len=1430 [TCP segment of a reassembled PDU]
22	0.142500	172.217.10.36	192.168.1.151	TCP	1116 HTTP/1.1 200 OK [TCP segment of a reassembled PDU]
23	0.142766	192.168.1.151	172.217.10.36	TCP	54 6495 → 80 [ACK] Seq=142 Ack=13933 Win=262144 Len=0
24	0.145744	172.217.10.36	192.168.1.151	HTTP	916 HTTP/1.1 200 OK (text/html)
25	0.167897	192.168.1.151	172.217.6.206	TCP	54 6494 → 80 [FIN, ACK] Seq=138 Ack=529 Win=262400 Len=0
26	0.170988	192.168.1.151	172.217.10.36	TCP	54 6495 → 80 [FIN, ACK] Seq=142 Ack=14795 Win=261120 Len=0
27	0.172374	172.217.6.206	192.168.1.151	TCP	60 80 → 6494 [FIN, ACK] Seq=529 Ack=139 Win=66816 Len=0
28	0.172660	192.168.1.151	172.217.6.206	TCP	54 6494 → 80 [ACK] Seq=139 Ack=530 Win=262400 Len=0
29	0.174762	172.217.10.36	192.168.1.151	TCP	60 80 → 6495 [FIN, ACK] Seq=14795 Ack=143 Win=61952 Len=0
30	0.175180	192.168.1.151	172.217.10.36	TCP	54 6495 → 80 [ACK] Seq=143 Ack=14796 Win=261120 Len=0

Step 2: Inspect the Trace

Select a packet for which the Protocol column is “HTTP” and the Info column says it is a GET. It is the packet that carries the web (HTTP) request sent from your computer to the server. (You can click the column headings to sort by that value, though it should not be difficult to find an HTTP packet by inspection.) Let’s have a closer look to see how the packet structure reflects the protocols that are in use.

1	0.000000	192.168.1.151	172.217.9.238	TCP	66 50359 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 WS=256 SACK_PERM=1
2	0.004604	172.217.9.238	192.168.1.151	TCP	66 80 → 50359 [SYN, ACK] Seq=0 Ack=1 Win=60720 Len=0 MSS=1380 SACK_PERM=1 WS=256
3	0.004883	192.168.1.151	172.217.9.238	TCP	54 50359 → 80 [ACK] Seq=1 Ack=1 Win=262144 Len=0
4	0.030124	192.168.1.151	172.217.9.238	HTTP	191 GET / HTTP/1.1
5	0.035254	172.217.9.238	192.168.1.151	TCP	60 80 → 50359 [ACK] Seq=1 Ack=138 Win=61952 Len=0
6	0.044222	172.217.9.238	192.168.1.151	HTTP	582 HTTP/1.1 301 Moved Permanently (text/html)
7	0.073539	192.168.1.151	172.217.3.100	TCP	66 50360 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 WS=256 SACK_PERM=1
8	0.077854	172.217.3.100	192.168.1.151	TCP	66 80 → 50360 [SYN, ACK] Seq=0 Ack=1 Win=65535 Len=0 MSS=1430 SACK_PERM=1 WS=256

wget "google.com"

Since we are fetching a web page, we know that the protocol layers being used are as shown below. That is, HTTP is the application layer web protocol used to fetch URLs. Like many Internet applications, it runs on top of the TCP/IP transport and network layer protocols. The link and physical layer protocols

depend on your network, but are typically combined in the form of Ethernet (shown) if your computer is wired, or 802.11 (not shown) if your computer is wireless.

With the HTTP GET packet selected, look closely to see the similarities and differences between it and our protocol stack as described next. The protocol blocks are listed in the middle panel. You can expand each block (by clicking on the “+” expander or icon) to see its details.

- The first Wireshark block is “Frame”. This is not a protocol, it is a record that describes overall information about the packet, including when it was captured and how many bits long it is.

```
▼ Frame 4: 191 bytes on wire (1528 bits), 191 bytes captured (1528 bits) on interface \Device\NPF_{FF8BDE0E-2083-4888-BFC9-2460E1FE7C6F}, id 0
  > Interface id: 0 (\Device\NPF_{FF8BDE0E-2083-4888-BFC9-2460E1FE7C6F})
    Encapsulation type: Ethernet (1)
    Arrival Time: Sep 30, 2020 16:33:02.722608000 Eastern Summer Time
    [Time shift for this packet: 0.000000000 seconds]
    Epoch Time: 1601497982.722608000 seconds
    [Time delta from previous captured frame: 0.025241000 seconds]
    [Time delta from previous displayed frame: 0.025241000 seconds]
    [Time since reference or first frame: 0.030124000 seconds]
    Frame Number: 4
    Frame Length: 191 bytes (1528 bits)
    Capture Length: 191 bytes (1528 bits)
    [Frame is marked: False]
    [Frame is ignored: False]
    [Protocols in frame: eth:ethertype:ip:tcp:http]
    [Coloring Rule Name: HTTP]
    [Coloring Rule String: http || tcp.port == 80 || http2]
```

- The second block is “Ethernet”. This matches our diagram! Note that you may have taken a trace on a computer using 802.11 yet still see an Ethernet block instead of an 802.11 block. Why? It happens because we asked Wireshark to capture traffic in Ethernet format on the cap-ture options, so it converted the real 802.11 header into a pseudo-Ethernet header.

```
▼ Ethernet II, Src: HewlettP_db:3e:fd (c8:d3:ff:db:3e:fd), Dst: Verizon_12:f1:29 (18:78:d4:12:f1:29)
  ▼ Destination: Verizon_12:f1:29 (18:78:d4:12:f1:29)
    Address: Verizon_12:f1:29 (18:78:d4:12:f1:29)
    .... ..0. .... = LG bit: Globally unique address (factory default)
    .... ..0. .... = IG bit: Individual address (unicast)
  ▼ Source: HewlettP_db:3e:fd (c8:d3:ff:db:3e:fd)
    Address: HewlettP_db:3e:fd (c8:d3:ff:db:3e:fd)
    .... ..0. .... = LG bit: Globally unique address (factory default)
    .... ..0. .... = IG bit: Individual address (unicast)
  Type: IPv4 (0x0800)
```


payload. You may have questions about the fields in each protocol as you look at them. We will explore these protocols and fields in detail in future labs.

Turn-in: Hand in your packet drawing.

HTTP	TCP	IP	Ethernet	
191 bytes	20 bytes	20 bytes	14 bytes	Get Packet

Step 4: Protocol Overhead

Estimate the download protocol overhead, or percentage of the download bytes taken up by protocol overhead. To do this, consider HTTP data (headers and message) to be useful data for the network to carry, and lower layer headers (TCP, IP, and Ethernet) to be the overhead. We would like this overhead to be small, so that most bits are used to carry content that applications care about. To work this out, first look at only the packets in the download direction for a single web fetch. You might sort on the Destination column to find them. The packets should start with a short TCP packet described as a SYN ACK, which is the beginning of a connection. They will be followed by mostly longer packets in the middle (of roughly 1 to 1.5KB), of which the last one is an HTTP packet. This is the main portion of the download. And they will likely end with a short TCP packet that is part of ending the connection. For each packet, you can inspect how much overhead it has in the form of Ethernet / IP / TCP headers, and how much useful HTTP data it carries in the TCP payload. You may also look at the HTTP packet in Wireshark to learn how much data is in the TCP payloads over all download packets.

Turn-in: Your estimate of download protocol overhead as defined above. Tell us whether you find this overhead to be significant.

Estimation:

1	0.000000	192.168.1.151	172.217.9.238	TCP	66 50359 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 WS=256 SACK_PERM=1
2	0.004604	172.217.9.238	192.168.1.151	TCP	66 80 → 50359 [SYN, ACK] Seq=0 Ack=1 Win=60720 Len=0 MSS=1380 SACK_PERM=1 WS=256
3	0.004883	192.168.1.151	172.217.9.238	TCP	54 50359 → 80 [ACK] Seq=1 Ack=1 Win=262144 Len=0
4	0.030124	192.168.1.151	172.217.9.238	HTTP	191 GET / HTTP/1.1
5	0.035254	172.217.9.238	192.168.1.151	TCP	60 80 → 50359 [ACK] Seq=1 Ack=138 Win=61952 Len=0
6	0.044222	172.217.9.238	192.168.1.151	HTTP	582 HTTP/1.1 301 Moved Permanently (text/html)
7	0.073539	192.168.1.151	172.217.3.100	TCP	66 50360 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 WS=256 SACK_PERM=1
8	0.077854	172.217.3.100	192.168.1.151	TCP	66 80 → 50360 [SYN, ACK] Seq=0 Ack=1 Win=65535 Len=0 MSS=1430 SACK_PERM=1 WS=256
9	0.077916	192.168.1.151	172.217.3.100	TCP	54 50360 → 80 [ACK] Seq=1 Ack=1 Win=262912 Len=0
10	0.099184	192.168.1.151	172.217.9.238	TCP	54 50359 → 80 [ACK] Seq=138 Ack=529 Win=261632 Len=0
11	0.101689	192.168.1.151	172.217.3.100	HTTP	195 GET / HTTP/1.1
12	0.106188	172.217.3.100	192.168.1.151	TCP	60 80 → 50360 [ACK] Seq=1 Ack=142 Win=66816 Len=0
13	0.163518	172.217.3.100	192.168.1.151	TCP	1484 80 → 50360 [ACK] Seq=1 Ack=142 Win=66816 Len=1430 [TCP segment of a reassembled PDU]
14	0.163518	172.217.3.100	192.168.1.151	TCP	1484 80 → 50360 [ACK] Seq=1431 Ack=142 Win=66816 Len=1430 [TCP segment of a reassembled PDU]
15	0.163518	172.217.3.100	192.168.1.151	TCP	1484 80 → 50360 [ACK] Seq=2861 Ack=142 Win=66816 Len=1430 [TCP segment of a reassembled PDU]
16	0.163518	172.217.3.100	192.168.1.151	TCP	1484 80 → 50360 [ACK] Seq=4291 Ack=142 Win=66816 Len=1430 [TCP segment of a reassembled PDU]
17	0.163518	172.217.3.100	192.168.1.151	TCP	1484 80 → 50360 [ACK] Seq=5721 Ack=142 Win=66816 Len=1430 [TCP segment of a reassembled PDU]
18	0.163518	172.217.3.100	192.168.1.151	TCP	1484 80 → 50360 [ACK] Seq=7151 Ack=142 Win=66816 Len=1430 [TCP segment of a reassembled PDU]
19	0.163518	172.217.3.100	192.168.1.151	TCP	1484 80 → 50360 [ACK] Seq=8581 Ack=142 Win=66816 Len=1430 [TCP segment of a reassembled PDU]
20	0.163518	172.217.3.100	192.168.1.151	TCP	1484 80 → 50360 [ACK] Seq=10011 Ack=142 Win=66816 Len=1430 [TCP segment of a reassembled PDU]
21	0.163518	172.217.3.100	192.168.1.151	TCP	1484 80 → 50360 [ACK] Seq=11441 Ack=142 Win=66816 Len=1430 [TCP segment of a reassembled PDU]
22	0.163518	172.217.3.100	192.168.1.151	TCP	1482 80 → 50360 [PSH, ACK] Seq=12871 Ack=142 Win=66816 Len=1428 [TCP segment of a reassembled PDU]
23	0.163972	192.168.1.151	172.217.3.100	TCP	54 50360 → 80 [ACK] Seq=142 Ack=14299 Win=262912 Len=0
24	0.168786	172.217.3.100	192.168.1.151	HTTP	493 HTTP/1.1 200 OK (text/html)
25	0.208890	192.168.1.151	172.217.3.100	TCP	54 50360 → 80 [ACK] Seq=142 Ack=14738 Win=262656 Len=0
26	0.221061	192.168.1.151	172.217.9.238	TCP	54 50359 → 80 [FIN, ACK] Seq=138 Ack=529 Win=261632 Len=0
27	0.225233	172.217.9.238	192.168.1.151	TCP	60 80 → 50359 [FIN, ACK] Seq=529 Ack=139 Win=61952 Len=0
28	0.225236	192.168.1.151	172.217.3.100	TCP	54 50360 → 80 [FIN, ACK] Seq=142 Ack=14738 Win=262656 Len=0
29	0.225537	192.168.1.151	172.217.9.238	TCP	54 50359 → 80 [ACK] Seq=139 Ack=530 Win=261632 Len=0
30	0.230040	172.217.3.100	192.168.1.151	TCP	60 80 → 50360 [FIN, ACK] Seq=14738 Ack=143 Win=66816 Len=0
31	0.230396	192.168.1.151	172.217.3.100	TCP	54 50360 → 80 [ACK] Seq=143 Ack=14739 Win=262656 Len=0

Frame #:	TCP IP Ethernet (Overhead) bytes	HTTP Data (Useful) bytes
2	66	0
3	54	0
4	40	191
5	60	0
8	66	0
9	54	0
10	54	0
11	40	195
12	60	0
13	1484	0
14	1484	0
15	1484	0
16	1484	0
17	1484	0
18	1484	0
19	1484	0

20	1484	0
21	1484	0
22	1482	0
23	54	0
24	40	493
25	54	0

Total Overhead: 15480 bytes

Useful data: 879 bytes

15480 bytes + 879 bytes = 16359 bytes

15480/16359 *100 = 94.63%

The overhead is significant as ~95% of the byte data consists of overhead. 5% is relevant (http) data.

Step 5: Demultiplexing Keys

When an Ethernet frame arrives at a computer, the Ethernet layer must hand the packet that it contains to the next higher layer to be processed. The act of finding the right higher layer to process received packets is called demultiplexing. We know that in our case the higher layer is IP. But how does the Ethernet protocol know this? After all, the higher-layer could have been another protocol entirely (such as ARP). We have the same issue at the IP layer – IP must be able to determine that the contents of IP message is a TCP packet so that it can hand it to the TCP protocol to process. The answer is that protocols use information in their header known as a “demultiplexing key” to determine the higher layer. CN5E Labs (1.0) © 2012 D. Wetherall 8

Turn-in: Hand in your answers to the [below] questions.

Look at the Ethernet and IP headers of a download packet in detail to answer the following questions:

1. Which Ethernet header field is the demultiplexing key that tells it the next higher layer is IP? What value is used in this field to indicate “IP”?


```

> Frame 13: 1484 bytes on wire (11872 bits), 1484 bytes captured (11872 bits) on interface \Device\NPF_{FF8BDE0E-2083-4888-BFC9-2460E1FE7C6F}, id 0
▼ Ethernet II, Src: Verizon_12:f1:29 (18:78:d4:12:f1:29), Dst: HewlettP_db:3e:fd (c8:d3:ff:db:3e:fd)
  ▼ Destination: HewlettP_db:3e:fd (c8:d3:ff:db:3e:fd)
    Address: HewlettP_db:3e:fd (c8:d3:ff:db:3e:fd)
    ....0. .... = LG bit: Globally unique address (factory default)
    ....0. .... = IG bit: Individual address (unicast)
  ▼ Source: Verizon_12:f1:29 (18:78:d4:12:f1:29)
    Address: Verizon_12:f1:29 (18:78:d4:12:f1:29)
    ....0. .... = LG bit: Globally unique address (factory default)
    ....0. .... = IG bit: Individual address (unicast)
  Type: IPv4 (0x0800)
> Internet Protocol Version 4, Src: 172.217.3.100, Dst: 192.168.1.151
> Transmission Control Protocol, Src Port: 80, Dst Port: 50360, Seq: 1, Ack: 142, Len: 1430

```

0000	c8 d3 ff db 3e fd 18 78 d4 12 f1 29 08 00 45 80>..x ...E.
0010	05 be f0 4a 00 00 3d 06 14 f3 ac d9 03 64 c0 a8	...J...=.....d..
0020	01 97 00 50 c4 b8 32 dd f5 c1 34 90 53 2c 50 10	...P...2...4..S,P.
0030	01 05 da 2d 00 00 48 54 54 50 2f 31 2e 31 20 32HT TP/1.1.2
0040	30 30 20 4f 4b 0d 0a 44 61 74 65 3a 20 57 65 64	00 OK..D ate: Wed
0050	2c 20 33 30 20 53 65 70 20 32 30 32 30 20 32 30	, 30 Sep 2020 20
0060	3a 33 33 3a 30 33 20 47 4d 54 0d 0a 45 78 70 69	:33:03 G MT..Expi

The protocol field is the demultiplexing key that tells it the next higher layer is IP. [08 00] is the value used to indicate IP.

2. Which IP header field is the demultiplexing key that tells it the next higher layer is TCP? What value is used in this field to indicate “TCP”?

Wireshark · Packet 13 · Ethernet (tcp port 80)

```

▼ Source: Verizon_12:f1:29 (18:78:d4:12:f1:29)
  Address: Verizon_12:f1:29 (18:78:d4:12:f1:29)
  ....0. .... = LG bit: Globally unique address (factory default)
  ....0. .... = IG bit: Individual address (unicast)
  Type: IPv4 (0x0800)
▼ Internet Protocol Version 4, Src: 172.217.3.100, Dst: 192.168.1.151
  0100 .... = Version: 4
  .... 0101 = Header Length: 20 bytes (5)
  > Differentiated Services Field: 0x80 (DSCP: CS4, ECN: Not-ECT)
  Total Length: 1470
  Identification: 0xf04a (61514)
  > Flags: 0x0000
  Fragment offset: 0
  Time to live: 61
  Protocol: TCP (6)
  Header checksum: 0x14f3 [validation disabled]
  [Header checksum status: Unverified]
  Source: 172.217.3.100
  Destination: 192.168.1.151

```

0010	05 be f0 4a 00 00 3d 06 14 f3 ac d9 03 64 c0 a8	...J...=.....d..
0020	01 97 00 50 c4 b8 32 dd f5 c1 34 90 53 2c 50 10	...P...2...4..S,P.

The protocol field is the demultiplexing key that tells it the next higher layer is IP. [06] is the value used to indicate IP.